

PSTAT174 Final Project 2

Week 9

Ryan Yee

5/29/2023

Contents

```
#install.packages("devtools")  
#devtools::install_github("FinYang/tsdl")  
#install.packages("astsa")
```

```
library(tsdl)  
library(astsa)
```

Source: U.S. Bureau of Economic Analysis

Release: Gross Domestic Product

Units: Billions of Chained 2012 Dollars, Not Seasonally Adjusted

Frequency: Quarterly

BEA Account Code: ND000334

U.S. Bureau of Economic Analysis, Real Gross Domestic Product [ND000334Q], retrieved from FRED, Federal Reserve Bank of St. Louis; <https://fred.stlouisfed.org/series/ND000334Q>, May 30, 2023.

```
gdp.csv = read.table("ND000334Q.csv", sep=",", header=FALSE, skip=1, nrows=85)  
head(gdp.csv)
```

```
##           V1           V2  
## 1 2002-01-01 3263.869  
## 2 2002-04-01 3362.508  
## 3 2002-07-01 3401.820  
## 4 2002-10-01 3460.159  
## 5 2003-01-01 3340.163  
## 6 2003-04-01 3429.079
```

```
tail(gdp.csv)
```


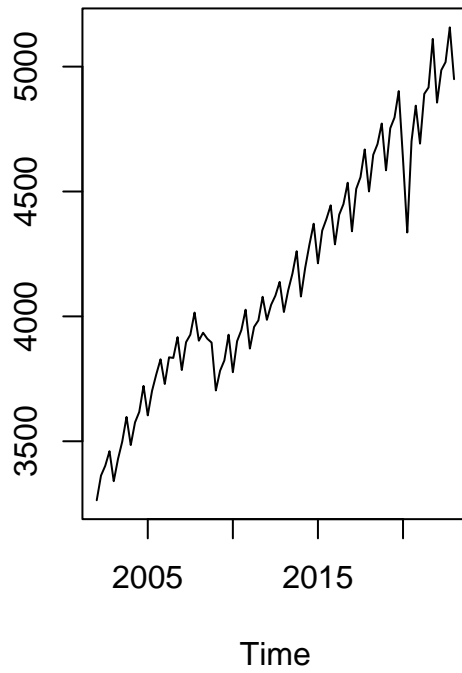
```
##           V1           V2
## 80 2021-10-01 5110.951
## 81 2022-01-01 4855.857
## 82 2022-04-01 4985.795
## 83 2022-07-01 5018.093
## 84 2022-10-01 5157.178
## 85 2023-01-01 4949.655
```

DISCLAIMER: The following data set is not gaussian distributed due to volatility during the Great Recession (December 2007 – June 2009) and the Pandemic (April 2020). The forecasting model does not pass the Shapiro-Wilk test!

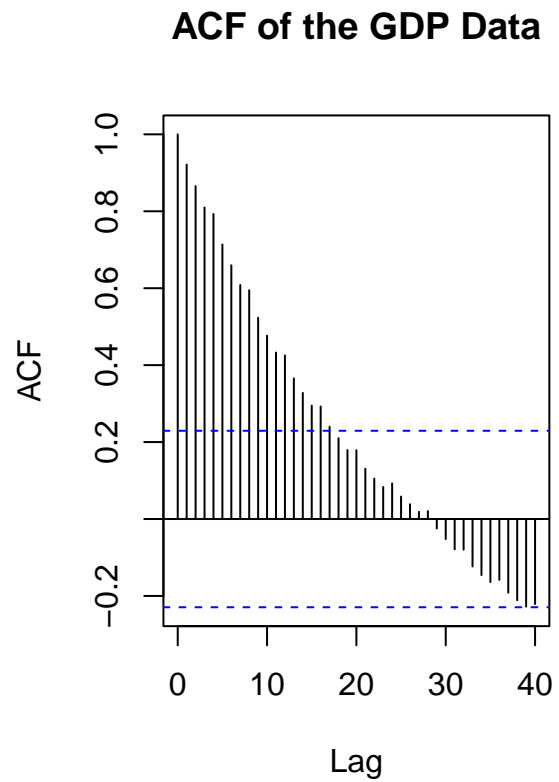
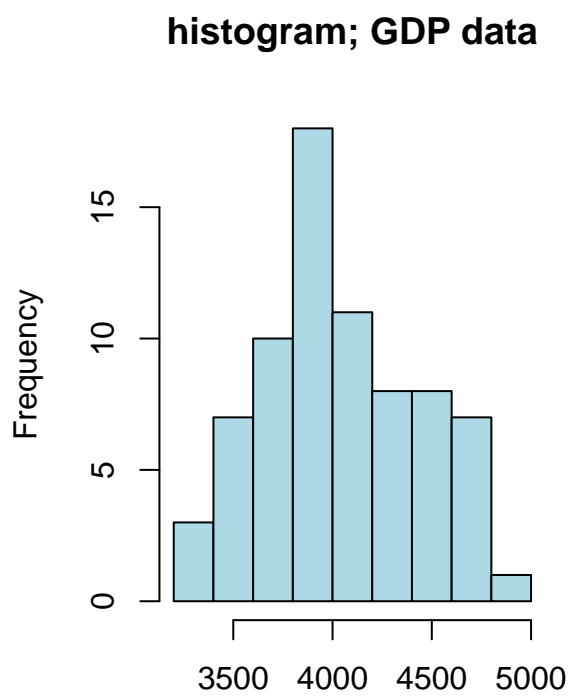
```
gdp = ts(gdp.csv[,2], start = c(2002,1), frequency = 4)
train <- gdp[1:73]
test <- gdp[74:85]
par(mfrow=c(1,2))
ts.plot(gdp, ylab = "Quarterly U.S. GDP (Billions of Chained 2012 Dollars)")
plot.ts(train)

fit <- lm(train ~ as.numeric(1:length(train))); abline(fit, col="red")
abline(h=mean(train), col="blue")
```

Quarterly U.S. GDP (Billions of Chained 2012 Dollars)




```
hist(train, col="light blue", xlab="", main="histogram; GDP data")
acf(train, lag.max=40, main="ACF of the GDP Data")
```



Histogram appear to be symmetric or bell-shaped (However, it is not gaussian). One might argue it is a little bit skewed right. Acf remains large.

Compare transformations





```
library(MASS)
bcTransform <- boxcox(train~ as.numeric(1:length(train)))
```




```
bcTransform$x[which(bcTransform$y == max(bcTransform$y))]
```

```
## [1] 0.1414141
```

```
lambda=bcTransform$x[which(bcTransform$y == max(bcTransform$y))]  
train.bc = (1/lambda)*(train^lambda-1)  
test.bc = (1/lambda)*(test^lambda-1)  
train.log <- log(train)  
train.sqrt = sqrt(train)  
  
op= par(mfrow=c(2,2))  
hist(train, col="light blue", xlab="", main="histogram; U_t")  
hist(train.log, col="light blue", xlab="", main="histogram; ln(U_t)")  
hist(train.bc, col="light blue", xlab="", main="histogram; bc(U_t)")  
hist(train.sqrt, col="light blue", xlab="", main="histogram; sqrt(U_t)")
```



```
#Compare transforms
op= par(mfrow=c(2,2))
ts.plot(train, main = "Original Times Series")
ts.plot(train.bc, main = "Box-Cox Transform")
ts.plot(train.log, main = "Log Transform")
ts.plot(train.sqrt, main = "Square Root Transform")
```




There isn't much difference between transformations. One might argue that a transformation is unnecessary. Choose $\ln(U_t)$ because it is more symmetric and seems to have more even variance.

Decomposition of $\ln(U_t)$ shows seasonality and almost linear trend

```
#install.packages("ggplot2")
#install.packages("ggfortify")
library(ggplot2)
library(ggfortify)

y <- ts(as.ts(train.log), frequency = 4)
decomp <- decompose(y)
plot(decomp)
```

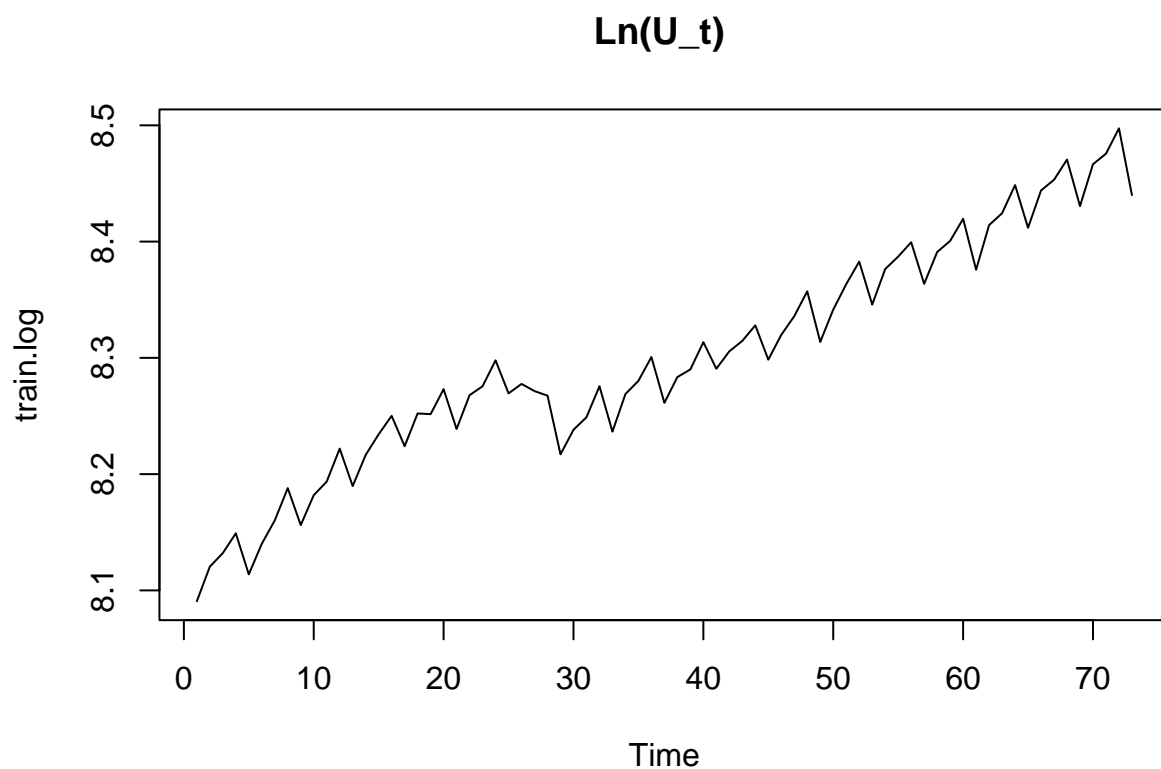
Decomposition of additive time series



```
var(train.log)
```

```
## [1] 0.009664501
```

```
plot.ts(train.log, main="Ln(U_t) ")
```

```
train.log_4 <- diff(train.log, lag=4)
plot.ts(train.log_4, main="Ln(U_t) differenced at lag 4")
var(train.log_4)
```


```
## [1] 0.0002872528
```

```
fit <- lm(train.log_4 ~ as.numeric(1:length(train.log_4))); abline(fit, col="red")
mean(train.log_4)
```

```
## [1] 0.02010507
```

```
abline(h=mean(train.log_4), col="blue")
```

Ln(U_t) differenced at lag 4



```
train.log_4_1 <- diff(train.log_4, lag=1)
plot.ts(train.log_4_1, main="Ln(Ut) differenced at lags 4 and then 1")
var(train.log_4_1)
```

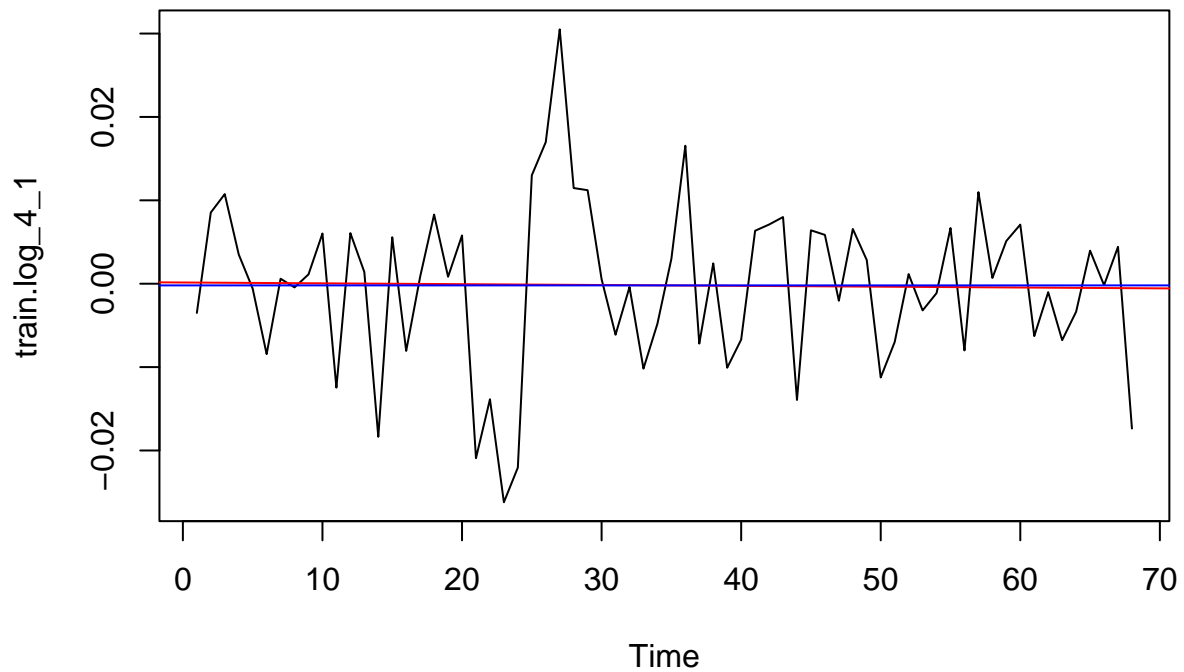
```
## [1] 9.788039e-05
```

```
fit <- lm(train.log_4_1 ~ as.numeric(1:length(train.log_4_1))); abline(fit, col="red")
mean(train.log_4_1)
```

```
## [1] -0.0002004654
```

```
abline(h=mean(train.log_4_1), col="blue")
```

Ln(U_t) differenced at lags 4 and then 1



Plot of $\ln(U_t)$

- Seasonality
- Trend
- Variance: 0.009665

Plot of $\ln(U_t)$ differenced at lag 4

- Seasonality no longer apparent
- Variance: 0.0002873 – lower!
- Trend is still here

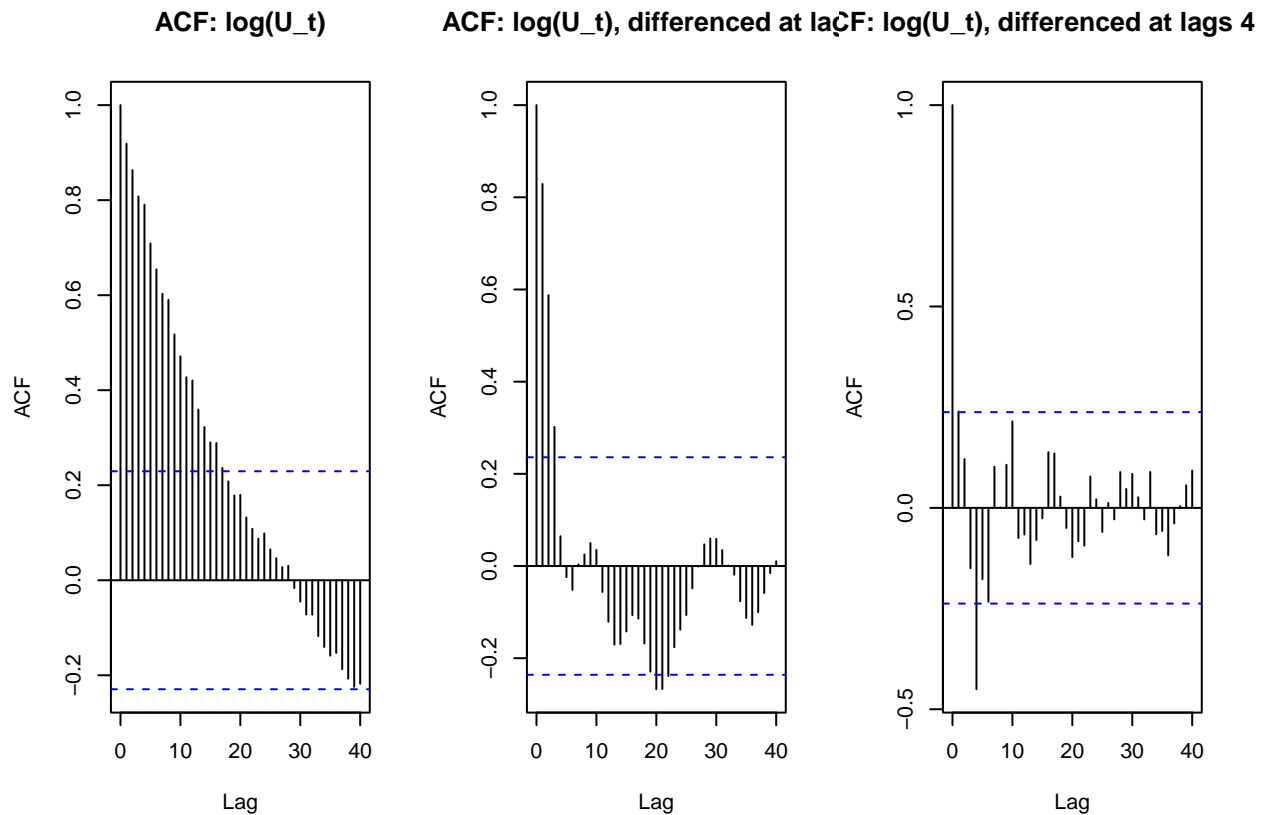
Plot of $\ln(U_t)$ differenced at lags 4 and then 1

- No Seasonality
- Variance: 9.788e-05 – even lower!
- No trend
- Data looks stationary, but check ACFs

```

par(mfrow=c(1, 3))
acf(train.log, lag.max=40, main="ACF: log(U_t)")
acf(train.log_4, lag.max=40, main="ACF: log(U_t), differenced at lag 4")
acf(train.log_4_1, lag.max=40, main="ACF: log(U_t), differenced at lags 4 and 1")

```



Plot of ACF of $\ln(U_t)$

- Slows decay indicated non-stationarity
- One sees seasonality

Plot of ACF of $\ln(U_t)$ differenced at lag 4

- Seasonality no longer apparent
- ACF decays slowly indicating non-stationarity

Plot of ACF of $\ln(U_t)$ differenced at lags 4 & 1

• ACF decay corresponds to a stationary process Conclude: Work with data $\ln(U_t)$ differenced at lags 4 & 1, U_t = the first 73 observations of the original data.


Histogram of $\nabla_1 \nabla_4 \ln(U_t)$ looks symmetric and almost Gaussian.

```

par(mfrow=c(1, 2))
hist(train.log_4_1, col="light blue", xlab="", main="histogram; ln(U_t) differenced at l
hist(train.log_4_1, density=20,breaks=20, col="blue", xlab="", main="Density of ln(U_t)
m<-mean(train.log_4_1)
std<- sqrt(var(train.log_4_1))
curve( dnorm(x,m,std), add=TRUE )

```

istogram; ln(U_t) differenced at lagsensity of ln(U_t) differenced at lags





ACF and PACF of $\text{train.log_4_1} = \nabla^1 \nabla^4 \ln(U_t)$,

```

par(mfrow=c(1, 2))
acf(train.log_4_1, lag.max=40, main="ACF of the log(U_t), differenced at lags 4 and 1")
pacf(train.log_4_1, lag.max=40, main="PACF of the ln(U_t), differenced at lags 4 and 1")

```

ACF of the $\log(U_t)$, differenced at lags 1 and 4



Determine possible candidate models $SARIMA(p, d, q) \times (P, D, Q)_s$ for the series $\ln(U_t)$.

Modeling the seasonal part (P, D, Q): For this part, focus on the seasonal lags $h = 1s, 2s$, etc.

- We applied one seasonal differencing so $D = 1$ at lag $s = 4$.
- The ACF shows a strong peak at $h = 1s$ and $h = 4s$.

A good choice for the MA part could be $Q = 1$ or $Q = 4$

- The PACF shows a peak at $h = 1s, 4s, 8s$ and $12s$. A good choice for the AR part could be $P = 1$ or $P = 4$.

Modeling the non-seasonal part (p, d, q): In this case focus on the within season lags, $h = 1, \dots, 11$.

- We applied one differencing to remove the trend: $d = 1$
- The ACF seems to be tailing off. Or perhaps cuts off at lag 1 or 4.

A good choice for the MA part could be $q = 1, 4$.

- The PACF cuts off at lag $h=1$ or 4.

A good choice for the AR part could be $p = 1, 4$.

As an illustration we fit the following model:

SARIMA $(p = 1, d = 1, q = 1) \times (P = 1, D = 1, Q = 1)_s=4$

SARIMA $(p = 4, d = 1, q = 4) \times (P = 4, D = 1, Q = 4)_s=4$

Try candidate models.

```
arima(train.log, order=c(4,1,4), seasonal = list(order = c(4,1,4), period = 4), method="ML")
```

```
## Warning in arima(train.log, order = c(4, 1, 4), seasonal = list(order = c(4, :  
## possible convergence problem: optim gave code = 1
```

```
##
```

```
## Call:
```

```
## arima(x = train.log, order = c(4, 1, 4), seasonal = list(order = c(4, 1, 4),  
##     period = 4), method = "ML")
```

```
##
```

```
## Coefficients:
```

```
## Warning in sqrt(diag(x$var.coef)): NaNs produced
```

```
##          ar1      ar2      ar3      ar4      ma1      ma2      ma3      ma4  
##      -0.1864  0.5091  0.0887 -0.8158  0.4534 -0.3167 -0.2387  0.4697  
## s.e.      NaN  0.1544  0.1215  0.0829      NaN      NaN  0.2366      NaN  
##          sar1      sar2      sar3      sar4      sma1      sma2      sma3      sma4  
##      -0.4179 -0.2166  0.3249  0.09  0.0358 -0.342 -0.7007  0.2117  
## s.e.      NaN      NaN      NaN      NaN      NaN      NaN  0.1359      NaN  
##
```

```
## sigma^2 estimated as 4.667e-05:  log likelihood = 237.54,  aic = -441.08
```

```
arima(train.log, order=c(1,1,1), seasonal = list(order = c(1,1,1), period = 4), method="ML")
```

```
##
```

```
## Call:
```

```
## arima(x = train.log, order = c(1, 1, 1), seasonal = list(order = c(1, 1, 1),  
##     period = 4), method = "ML")
```

```
##
```

```
## Coefficients:
```

```
##          ar1          ma1      sar1      sma1
```

```

##          0.4739 -0.2655  0.0334 -0.8102
## s.e.    0.3534   0.3747  0.1762   0.1477
##
## sigma^2 estimated as 5.746e-05:  log likelihood = 233.48,  aic = -456.96

arima(train.log, order=c(1,1,1), seasonal = list(order = c(1,1,1), period = 4), fixed =

## Warning in arima(train.log, order = c(1, 1, 1), seasonal = list(order = c(1, :
## some AR parameters were fixed: setting transform.pars = FALSE

##
## Call:
## arima(x = train.log, order = c(1, 1, 1), seasonal = list(order = c(1, 1, 1),
##      period = 4), fixed = c(NA, 0, 0, NA), method = "ML")
##
## Coefficients:
##          ar1  ma1  sar1      sma1
##          0.2132    0    0 -0.7857
## s.e.    0.1246    0    0  0.1197
##
## sigma^2 estimated as 5.808e-05:  log likelihood = 233.2,  aic = -460.39

arima(train.log, order=c(1,1,0), seasonal = list(order = c(0,1,1), period = 4), method="

##
## Call:
## arima(x = train.log, order = c(1, 1, 0), seasonal = list(order = c(0, 1, 1),
##      period = 4), method = "ML")
##
## Coefficients:
##          ar1      sma1
##          0.2132 -0.7857
## s.e.    0.1246  0.1197
##
## sigma^2 estimated as 5.808e-05:  log likelihood = 233.2,  aic = -460.39

arima(train.log, order=c(0,1,4), seasonal = list(order = c(1,1,1), period = 4), method="

##
## Call:
## arima(x = train.log, order = c(0, 1, 4), seasonal = list(order = c(1, 1, 1),
##      period = 4), method = "ML")

```



```
##
## Coefficients:
##          ma1      ma2      ma3      ma4      sar1      sma1
##          0.1984  0.2131 -0.0709  0.8025 -0.6712 -0.7937
## s.e.    0.0990  0.0906   0.1167  0.1242   0.1412   0.1261
##
## sigma^2 estimated as 5.009e-05:  log likelihood = 236.7,  aic = -459.39

arima(train.log, order=c(0,1,4), seasonal = list(order = c(1,1,1), period = 4), fixed =

##
## Call:
## arima(x = train.log, order = c(0, 1, 4), seasonal = list(order = c(1, 1, 1),
##      period = 4), fixed = c(NA, NA, 0, NA, NA, NA), method = "ML")
##
## Coefficients:
##          ma1      ma2      ma3      ma4      sar1      sma1
##          0.1811  0.2353      0  0.8422 -0.6987 -0.7934
## s.e.    0.0826  0.0782      0  0.1030   0.1465   0.1237
##
## sigma^2 estimated as 5.1e-05:  log likelihood = 236.48,  aic = -460.96
```


Our chosen model is: SARMA ($p = 0, d = 1, q = 4$) \times ($P = 1, D = 1, Q = 1$) $s=4$

```
fit.i <- sarima(xdata = train.log, p = 0, d = 1, q = 4, P = 1, D = 1, Q = 1, S = 4)

## initial  value -4.609268
## iter    2 value -4.642214
## iter    3 value -4.800009
## iter    4 value -4.811586
## iter    5 value -4.819757
## iter    6 value -4.822514
## iter    7 value -4.824085
## iter    8 value -4.826922
## iter    9 value -4.829141
## iter   10 value -4.830885
## iter   11 value -4.834626
## iter   12 value -4.842839
## iter   13 value -4.850116
## iter   14 value -4.857194
## iter   15 value -4.861821
```

```
## iter 16 value -4.867761
## iter 17 value -4.872671
## iter 18 value -4.877943
## iter 19 value -4.883747
## iter 20 value -4.890798
## iter 21 value -4.893435
## iter 22 value -4.893997
## iter 23 value -4.896854
## iter 24 value -4.897684
## iter 25 value -4.906952
## iter 26 value -4.910366
## iter 27 value -4.932993
## iter 28 value -4.950032
## iter 29 value -4.963518
## iter 30 value -4.974003
## iter 31 value -4.982583
## iter 32 value -4.983232
## iter 33 value -4.987585
## iter 34 value -4.991273
## iter 35 value -5.005516
## iter 36 value -5.015382
## iter 37 value -5.026687
## iter 38 value -5.032704
## iter 39 value -5.041045
## iter 40 value -5.043318
## iter 41 value -5.047523
## iter 42 value -5.054274
## iter 43 value -5.057815
## iter 44 value -5.060731
## iter 45 value -5.060856
## iter 46 value -5.065235
## iter 47 value -5.067495
## iter 48 value -5.070624
## iter 49 value -5.071428
## iter 50 value -5.075305
## iter 51 value -5.078713
## iter 52 value -5.081303
## iter 53 value -5.084384
## iter 54 value -5.086204
## iter 55 value -5.089075
## iter 56 value -5.089132
## iter 57 value -5.089195
## iter 58 value -5.089196
## iter 59 value -5.090696
## iter 59 value -5.090696
```

```
## iter 60 value -5.090790
## iter 61 value -5.090792
## iter 62 value -5.091655
## iter 63 value -5.092347
## iter 64 value -5.092487
## iter 65 value -5.093596
## iter 66 value -5.093611
## iter 67 value -5.093669
## iter 68 value -5.093670
## iter 69 value -5.093670
## iter 69 value -5.093670
## iter 70 value -5.094122
## iter 71 value -5.094122
## iter 71 value -5.094122
## iter 72 value -5.094234
## iter 72 value -5.094234
## iter 73 value -5.094237
## iter 73 value -5.094237
## iter 74 value -5.094238
## iter 74 value -5.094238
## iter 75 value -5.094239
## iter 75 value -5.094239
## iter 75 value -5.094239
## final value -5.094239
## converged
## initial value -4.898013
## iter 2 value -4.899022
## iter 3 value -4.899519
## iter 4 value -4.899755
## iter 5 value -4.899773
## iter 6 value -4.899778
## iter 7 value -4.899779
## iter 8 value -4.899780
## iter 9 value -4.899780
## iter 9 value -4.899780
## iter 9 value -4.899780
## final value -4.899780
## converged
```



```
print('Coefficients')
```

```
## [1] "Coefficients"
```

```
fit.i$fit$coef
```

```
##          ma1          ma2          ma3          ma4          sar1          sma1
## 0.19838235 0.21305446 -0.07094011 0.80250441 -0.67129980 -0.79368654
```

```
polyroot(c(1,0.19838,0.21305,-0.07094,0.80250))
```

```
## [1] 0.7261190+0.8478433i -0.6819196+0.7314374i -0.6819196-0.7314374i
## [4] 0.7261190-0.8478433i
```

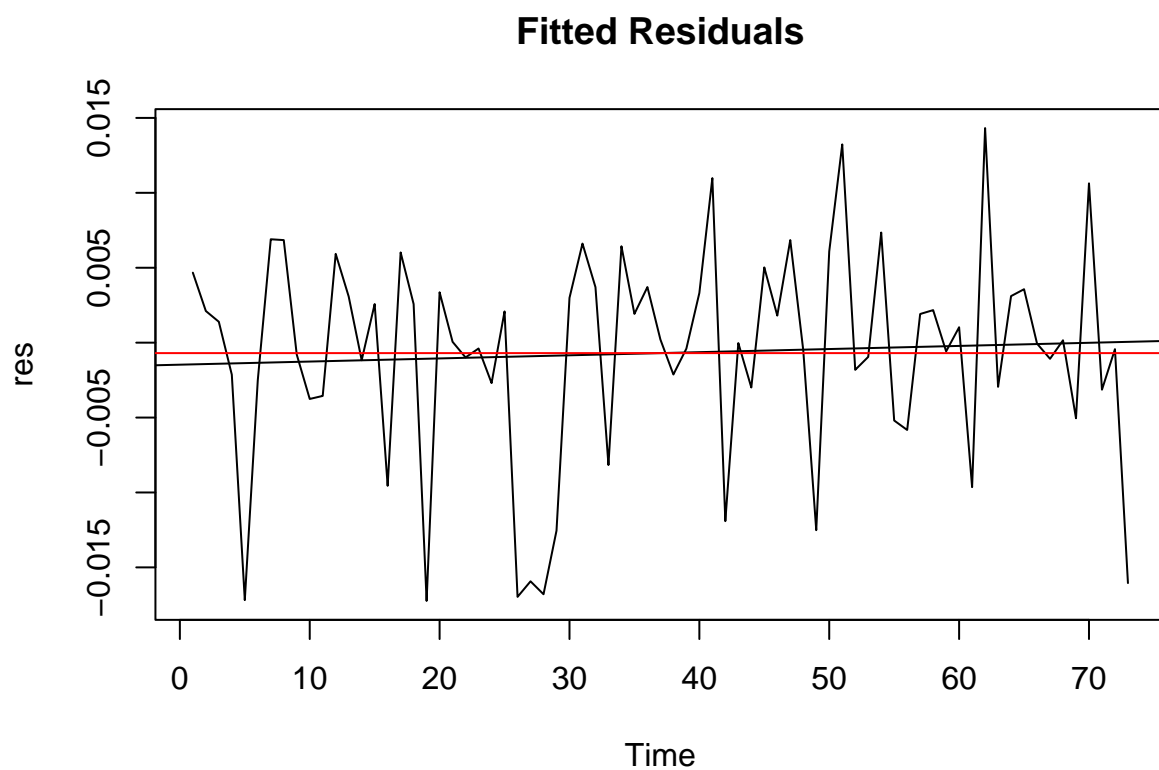
model is invertible.

```
# Residual plots:
res <- fit.i$fit$residuals
mean(res); var(res)
```

```
## [1] -0.0006974634
```

```
## [1] 5.137446e-05
```


```
# layout(matrix(c(1, 1, 2, 3), 2, 2, byrow=T))
par(mfrow=c(1, 1))
ts.plot(res, main="Fitted Residuals")
t <- 1:length(res)
fit.res = lm(res~ t)
abline(fit.res)
abline(h = mean(res), col = "red")
```



```
var(fit.res$residuals)
```

```
## [1] 5.117826e-05
```

```
# ACF and PACF:
par(mfrow=c(1, 2))
acf(res, main="Autocorrelation")
pacf(res, main="Partial Autocorrelation")
```



All acf and pacf of residuals are within confidence intervals and can be counted as zeros!

```
# Test for independence of residuals
Box.test(res, lag = 9, type = c("Box-Pierce"), fitdf = 6)
```

```
##
## Box-Pierce test
##
## data: res
## X-squared = 1.8362, df = 3, p-value = 0.6071
```

```
Box.test(res, lag = 9, type = c("Ljung-Box"), fitdf = 6)
```

```
##  
## Box-Ljung test  
##  
## data: res  
## X-squared = 2.0388, df = 3, p-value = 0.5644
```



```
Box.test(res^2, lag = 9, type = c("Ljung-Box"), fitdf = 0)
```

```
##  
## Box-Ljung test  
##  
## data: res^2  
## X-squared = 11.61, df = 9, p-value = 0.2362
```



```
# Test for normality of residuals  
shapiro.test(res)
```

```
##  
## Shapiro-Wilk normality test  
##  
## data: res  
## W = 0.93546, p-value = 0.001031
```

```
# Histogram and QQ-plot:  
par(mfrow=c(1,2))  
hist(res, main = "Histogram")  
qqnorm(res)  
qqline(res, col = "blue")
```



```
par(mfrow=c(1, 2))
hist(res, col="light blue", xlab="", main="histogram; res")
hist(res, density=20,breaks=20, col="blue", xlab="", main="Density of res", prob=TRUE)
m<-mean(res)
#0.001234
std<- sqrt(var(res))
curve( dnorm(x,m,std), add=TRUE )
```

```
# Access the coefficients of the SARIMA model
```

```
coefficients <- coef(fit.i$fit)
```

```
# Print the coefficients
```

```
coefficients
```

```
##          ma1          ma2          ma3          ma4          sar1          sma1
## 0.19838235 0.21305446 -0.07094011 0.80250441 -0.67129980 -0.79368654
```

```
# Access the variance
```

```
var <- fit.i$fit$sigma2
```

```
var
```

```
## [1] 5.009395e-05
```

The residuals appear to be white noise and they are normally distributed since the histogram appears approximately symmetric and the QQ plot shows residuals aligning along the diagonal line.


The SARIMA(0, 1, 0) \times (1, 1, 1)₁₂ model with coefficients (sar1 = -0.04274, sma1 = -0.40209) can be expressed as:

$$(1 - (-0.04274))(1 - B^{12})Y_t = (1 - (-0.40209)B)(1 - B)Z_t$$


Here, Y_t represents the series “train.log_4_1” and Z_t is a normally distributed error term. The coefficients sar1 and sma1 correspond to the autoregressive (AR) and seasonal moving average (SMA) terms, respectively, in the SARIMA model.

Hence, the final model will be $(1 + 0.04274)(1 - B^{12})Y_t = (1 + 0.40209)B(1 - B)Z_t$, where $Z_t \sim N(0, 1503)$.


```
x <- sarima.for(xdata=train.log, n.ahead = 12, p=3, d=1, q=2, P=0, D=0, Q=0, S=4)
points(1:length(train), train.log, col = "red", pch = 19, cex = 0.8)
points(length(train) + 1:length(test), log(test), col = "blue", pch = 19, cex = 0.8)
```




```
x <- sarima.for(xdata=train.log, n.ahead = 12, p=4, d=1, q=4, P=4, D=1, Q=4, S=4)
points(1:length(train), train.log, col = "red", pch = 19, cex = 0.8)
points(length(train) + 1:length(test), log(test), col = "blue", pch = 19, cex = 0.8)
```



```
x <- sarima.for(xdata=train.log, n.ahead = 12, p=0, d=1, q=0, P=0, D=1, Q=4, S=4)
points(1:length(train), train.log, col = "red", pch = 19, cex = 0.8)
points(length(train) + 1:length(test), log(test), col = "blue", pch = 19, cex = 0.8)
```



```
x <- sarima.for(xdata=train.log, n.ahead = 12, p=0, d=1, q=4, P=1, D=1, Q=1, S=4)
points(1:length(train), train.log, col = "red", pch = 19, cex = 0.8)
points(length(train) + 1:length(test), log(test), col = "blue", pch = 19, cex = 0.8)
```



```
#install.packages("forecast") # Install the forecast package if not already installed
library(forecast)              # Load the forecast package
```

```
## Registered S3 method overwritten by 'quantmod':
##   method      from
##   as.zoo.data.frame zoo
```


```
## Registered S3 methods overwritten by 'forecast':
##   method      from
##   autoplot.Arima      ggfortify
##   autoplot.acf        ggfortify
##   autoplot.ar         ggfortify
##   autoplot.bats       ggfortify
##   autoplot.decomposed.ts ggfortify
##   autoplot.ets        ggfortify
##   autoplot.forecast    ggfortify
##   autoplot.stl        ggfortify
##   autoplot.ts         ggfortify
```

```
## fitted.ar          ggfortify
## fortify.ts         ggfortify
## residuals.ar       ggfortify
```


```
##
## Attaching package: 'forecast'
```

```
## The following object is masked from 'package:astsa':
##
## gas
```

```
pred.tr <- sarima.for(xdata=train.log, n.ahead = 12, p=0, d=1, q=4, P=1, D=1, Q=1, S=4)
```




```
U.tr= pred.tr$pred + 2*pred.tr$se
L.tr= pred.tr$pred - 2*pred.tr$se
ts.plot(train.log, xlim=c(1,length(train.log)+12), ylim = c(min(train.log),max(U.tr)))
lines(U.tr, col="blue", lty="dashed")
lines(L.tr, col="blue", lty="dashed")
points((length(train.log)+1):(length(train.log)+12), pred.tr$pred, col="red")
points(length(train) + 1:length(test), log(test), col = "blue", pch = 19, cex = 0.8)
```




```

pred.orig <- exp(pred.tr$pred)
U= exp(U.tr)
L= exp(L.tr)
ts.plot(train, xlim=c(1,length(train)+12), ylim = c(min(train),max(U)))
lines(U, col="blue", lty="dashed")
lines(L, col="blue", lty="dashed")
points((length(train)+1):(length(train)+12), pred.orig, col="red")


```



```
ts.plot(gdp.csv, xlim=c(1,length(train)+12), ylim = c(min(train),max(U)), col="red")
lines(U, col="blue", lty="dashed")
lines(L, col="blue", lty="dashed")
points((length(train)+1):(length(train)+12), pred.orig, col="black")
```

```
ts.plot(train, xlim = c(65,length(train)+12), ylim = c(2000,max(U)))
lines(U, col="blue", lty="dashed")
lines(L, col="blue", lty="dashed")
points((length(train)+1):(length(train)+12), pred.orig, col="red")
```



```
ts.plot(gdp.csv, xlim = c(65,length(train)+12), ylim = c(2000,max(U)), col="red")
lines(U, col="blue", lty="dashed")
lines(L, col="blue", lty="dashed")
points((length(train)+1):(length(train)+12), pred.orig, col="black")
```

